




Saved 20-node graph, seed 0: best cost found within 500 iterations (3 network seeds each; identical values overlap)



The same runs: cost at iteration 500 (above the best = the run drifted away from it)



Phase 2: learned against uniform attention on the same run. 256 of 270 pairs reach exactly the same best cost; 270 of 270 violate the same number of constraints (the rest differ only by tie-break terms, at most 0.019)



Where the learned damping weight went in the 204 runs that learn it: every run stayed between 0.49962 and 0.50043, while the network could reach anything from 0.269 to 0.731 (inset)

